

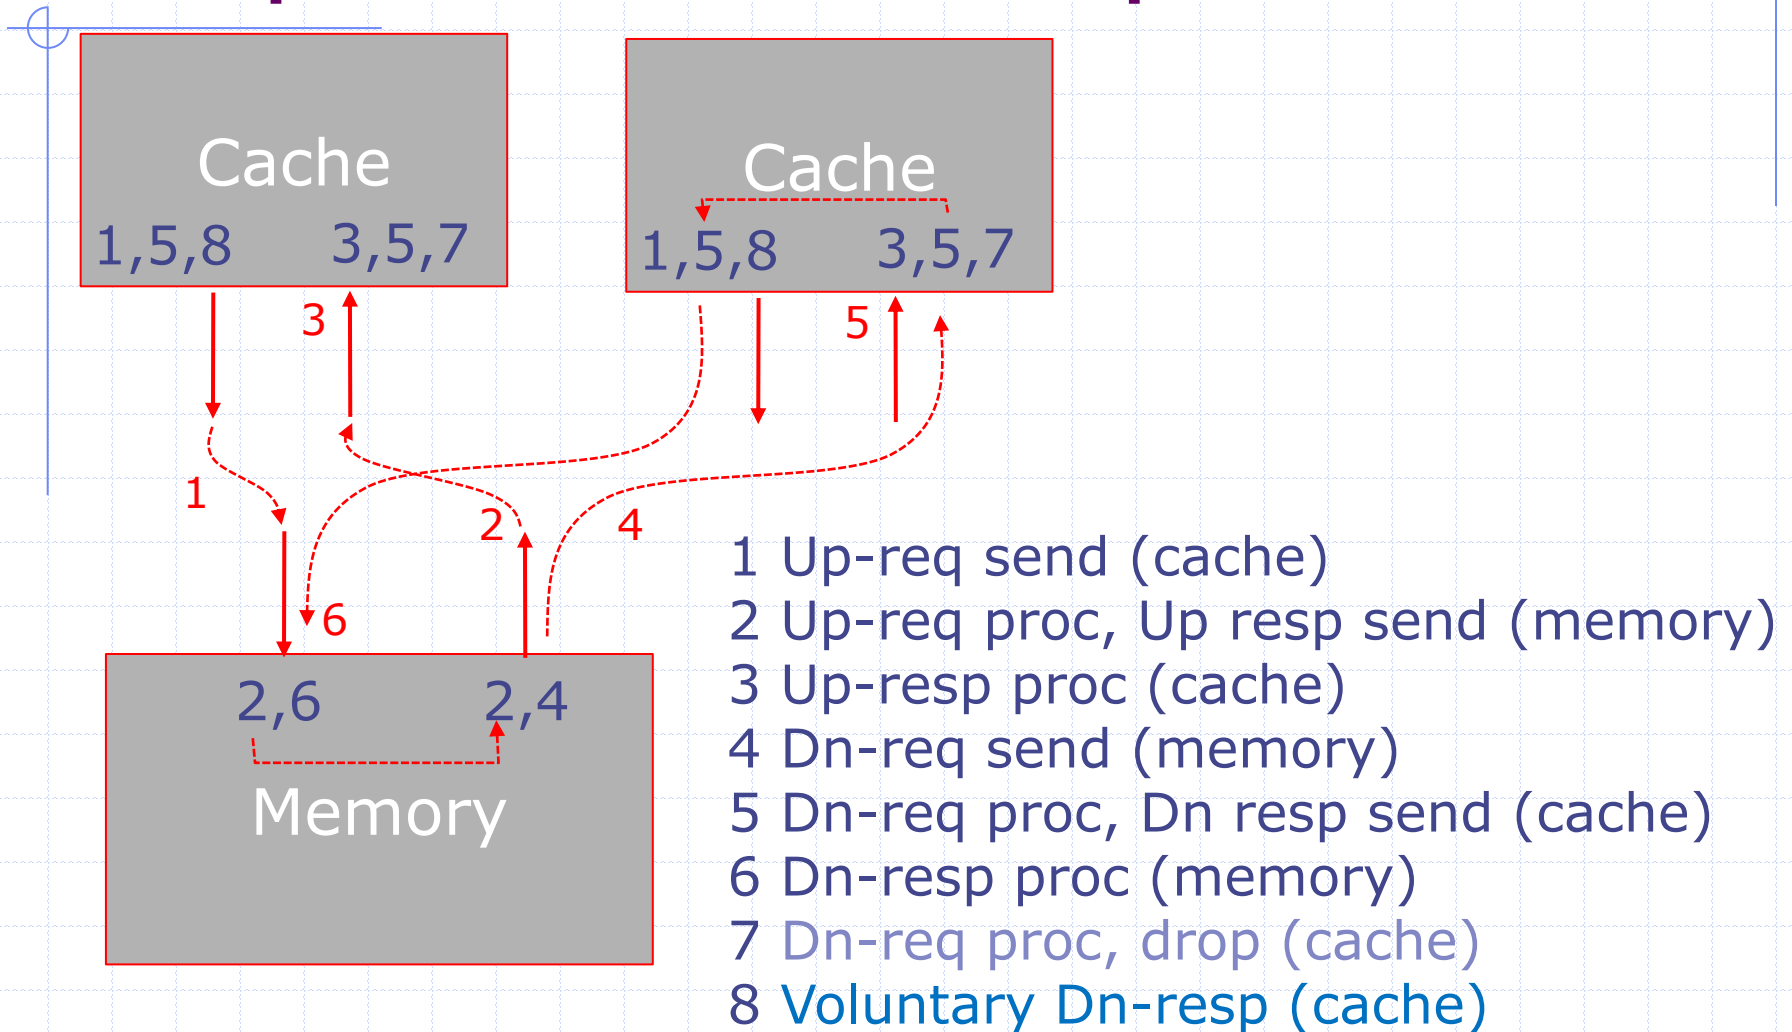
CC protocol for blocking caches

Extension to the cache rules
for Blocking L1 design
discussed in lecture L15

Code is somewhat simplified by assuming that
cache-line size = one word


Some pseudo syntax is used

Processing misses: Requests and Responses



Req method hit processing

```
method Action req(MemReq r) if(mshr == Ready);  
  let a = r.addr;  
  let hit = contains(state, a);  
  if(hit) begin  
    let slot = getSlot(state, a);  
    let x = dataArray[slot];  
    if(r.op == Ld) hitQ.enq(x);  
    else // it is store  
      if (isStateM(state[slot])  
        dataArray[slot] <= r.data;  
      else begin missReq <= r; mshr <= SendFillReq;  
        missSlot <= slot; end  
    end  
  else begin missReq <= r; mshr <= StartMiss; end // (1)  
endmethod
```



Start-miss and Send-fill rules

Rdy -> StrtMiss -> SndFillReq -> WaitFillResp -> Resp -> Rdy

```
rule startMiss(mshr == StartMiss);  
  let slot = findVictimSlot(state, missReq.addr);  
  if(!isStateI(state[slot]))  
    begin // write-back (Evacuate)  
      let a = getAddr(state[slot]);  
      let d = (isStateM(state[slot])? dataArray[slot]: -);  
      state[slot] <= (I, _);  
      c2m.enq(<Resp, c->m, a, I, d>); end  
  mshr <= SendFillReq; missSlot <= slot; endrule  
  
rule sendFillReq (mshr == SendFillReq);  
  // must have a reserved slot to receive reply  
  let upg = (missReq.op == Ld)? S : M;  
  c2m.enq(<Req, c->m, missReq.addr, upg, - >);  
  mshr <= WaitFillResp; endrule // (1)
```

Wait-fill rule and Proc Resp rule

Rdy -> StrtMiss -> SndFillReq -> WaitFillResp -> Resp -> Rdy

```
rule waitFillResp ((mshr == WaitFillResp) &&&
    (m2c.firstResp matches <Resp, m->c, .a, .cs, .d>));
    let slot = missSlot;
    dataArray[slot] <=
        (missReq.op == Ld)? d : missReq.data;
    state[slot] <= (cs, a);
    m2c.deqResp;
    mshr <= Resp;
endrule // (3)

rule sendProc(mshr == Resp);
    if (missReq.op == Ld) begin let slot = missSlot;
        c2p.enq(dataArray[slot]); end
    mshr <= Ready;
endrule
```

Parent Responds

```
rule parentResp
```

```
    (c2m.firstReq matches <Req, .c->m, .a, .y, .*>);
```

```
    let slot = getSlot(state, a); // in a 2-level
    // system a has to be present in the memory
```

```
    let statea = state[slot];
```

```
    if(( $\forall i \neq c$ , isCompatible(statea.dir[i], y))
```

```
        && (statea.waitc[c]=No)) begin
```

```
        let d =(statea.dir[c]=I)? dataArray[slot]: -);
```

```
        m2c.enq(<Resp, m->c, a, y, d>);
```

```
        state[slot].dir[c]:=y;
```

```
        c2m.deq;
```

```
    end
```

```
endrule
```

IsCompatible(M, M) = False
IsCompatible(M, S) = False
IsCompatible(S, M) = False
All other cases = True

Parent (Downgrade) Requests

rule dwn

(c2m.firstReq matches <Req, c->m, .a, .y, .*>);

let slot = getSlot(state, a);

let statea = state[slot];

if (findChild2Dwn(statea) matches (Valid .i))

begin

state[slot].waitc[i] <= Yes;

m2c.enq(<Req, m->i, a, (y==M?I:S), ? >);

end;

endrule // (4)

This rule will execute as long some child cache is not compatible with the incoming request

Parent receives Response

```
rule dwnRsp
```

```
(c2m.firstResp matches <Resp, c->m, .a, .y, .data>);  
c2m.deqResp;
```

```
let slot = getSlot(state, a);
```

```
let statea = state[slot];
```

```
if(statea.dir[c]=M) dataArray[slot]<=data;
```

```
state[slot].dir[c]<=y;
```

```
state[slot].waitc[c]<=No;
```

```
endrule // (6)
```


Child Responds

Incoming downgrade requests in L1

```
rule dng ((mshr != Resp) &&&
           (m2c.firstReq matches <Req,m->c,.a,.y,.*>));
let slot = getSlot(state,a);
if (getCacheState(state[slot])>y) begin
    let d = (isStateM(state[slot])? dataArray[slot]: -);
    c2m.enq(<Resp, c->m, a, y, d>);
    state[slot] <= (y,a);
end
// the address has already been downgraded
m2c.deqReq;
endrule // (5) and (7)
```

Child Voluntarily downgrades

```
rule startMiss(mshr == Ready);  
  let slot = findVictimSlot(state);  
  if(!isStateI(state[slot]))  
    begin // write-back (Evacuate)  
      let a = getAddr(state[slot]);  
      let d = (isStateM(state[slot])? dataArray[slot]: -);  
      state[slot] <= (I, _);  
      c2m.enq(<Resp, c->m, a, I, d>);  
    end  
endrule // (8)
```

Rules 1 to 8 are complete - cover all possibilities
and cannot deadlock or violate cache invariants

MSI protocol: some issues

- ◆ It never makes sense to have two outstanding requests for the same address from the same processor/cache
- ◆ It is possible to have multiple requests for the same address from different processors. Hence there is a need to arbitrate requests
- ◆ A cache needs to be able to evict an address in order to make room for a different address
 - Voluntary downgrade
- ◆ Memory system (higher-level cache) should be able to force a lower-level cache to downgrade
 - caches need to keep track of the state of their children's caches

Invariants for a CC-protocol design

- ◆ Directory state is always a conservative estimate of a child's state
 - E.g., if directory thinks that a child cache is in S state then the cache has to be in either I or S state
- ◆ For every request there is a corresponding response, though sometimes it is generated even before the request is processed
- ◆ Communication system has to ensure that
 - responses cannot be blocked by requests
 - a request cannot overtake a response for the same address
- ◆ At every merger point for requests, we will assume fair arbitration to avoid starvation

Nonblocking Synchronization

Load-reserve & Store-conditional

Special register(s) to hold reservation flag and address, and the outcome of store-conditional

Load-reserve R, m:

```
<flag, adr> ← <1, m>;  
R ← M[m];
```

Store-conditional m, R:

```
if <flag, adr> == <1, m>  
then cancel other procs'  
    reservation on m;  
    M[m] ← R;  
    status ← succeed;  
else status ← fail;
```

try:
spin:

```
Load-reserve Rhead, head  
Load Rtail, tail  
if Rhead == Rtail goto spin  
Load R, (Rhead)  
Rhead = Rhead + 1  
Store-conditional head, Rhead  
if (status == fail) goto try  
process(R)
```

The corresponding instructions in RISC V are called lr and sc, respectively

Nonblocking Synchronization

Load-reserve R, (m):
 <flag, adr> \leftarrow <1, m>;
 R \leftarrow M[m];

Store-conditional (m), R:
 if <flag, adr> == <1, m>
 then cancel other procs'
 reservation on m;
 M[m] \leftarrow R;
 status \leftarrow succeed;
 else status \leftarrow fail;

- ◆ The flag is cleared in other processors on a Store using the CC protocol's invalidation mechanism
- ◆ Usually address m is not remembered by Load-reserve; the flag is cleared on *any* invalidation
 - works as long as the Load-reserve instructions are not used in a nested manner
- ◆ These instructions won't work properly if Loads and Stores can be reordered dynamically